

1 Define the target

strong

Name the task and the sites; fix a worst-site acceptance level first.

Canine: mean AUROC .79 hid a scanner at .53 (ERM).

2 Collect sites, not slides

strong

Hold out whole sites; size the number of sites from a pilot.

A .05 gain needs 7–53 sites (tumor cohorts).

3 Train once, strongly

moderate

Fine-tune a pathology FM with color jitter; tune its own LR.

Jitter beat ERM in 18/20 domains; FM + jitter worst site .86–.98.

4 Adapt without labels

mixed

Optional: re-estimate BN statistics or normalize stain per slide.

BN-adapt +0.12 on canine (5/5); -0.004 on MIDOG++ (no gain).

5 Select and validate

strong

Select on training-site validation; report every site, seeds, ECE.

Selection rules agree at standard scale ($\tau_b \geq .83$).

6 Accept each new site

moderate

Label a small sample; run the site acceptance test; monitor.

False acceptance $\leq 2.1\%$ (prior alone 8.6%); 100 labels: 84% vs 70% power.

Evidence grades: strong = holds in all four cohorts at the standard tier; moderate = consistent direction but not significant after correction, or holding with exceptions; mixed = large gains in some cohorts, none in others. Numbers: standard tier, ID-val selection.